

A Multi-Stream Temporal Context Model for Narrative Memory

Xinming Xu and Jeremy R. Manning*

Department of Psychological and Brain Sciences, Dartmouth College

*Corresponding author: jeremy.r.manning@dartmouth.edu

Abstract

The Temporal Context Model¹⁰ and its extension the Context Maintenance and Retrieval model¹⁵ together provide a widely-used computational framework for understanding how memory for a sequence of events depends on the temporal context in which they were encoded. In both, a single temporal-context vector drifts continuously through encoding, so the similarity between the contexts of two events falls off smoothly with the number of intervening events. This single-stream assumption explains a broad range of phenomena in word-list free recall, but it faces a puzzle when memory for *naturalistic* narratives is studied. In pilot data from a cued-recall study of a television show whose plot interleaves three storylines, Xu et al.²⁰ find that cued recall of two narratively adjacent events is essentially unchanged by whether they are temporally adjacent (the *grouped* condition) or are separated by 2 to 7 intervening events from other storylines (the *bridge* condition), with a Bayes factor of $BF_{01} = 412$ for the null of no difference. Under a single-stream context model, the bridge condition should incur substantial context drift and therefore a substantial cued-recall deficit – the observed equivalence is unexpected. Here we propose a *Multi-Stream Temporal Context Model* (MS-TCM) that augments the single global temporal context with a bank of storyline-specific context vectors that drift only while their own storyline is active and are held constant (“frozen”) otherwise. We derive closed-form parameter regimes under which MS-TCM predicts the grouped-bridge equivalence, and fit the model by

maximum likelihood, with participant-level bootstrap confidence intervals, to the category condition of the feature-rich free-recall dataset of Manning et al.¹³ as a worked example. The MLE strongly prefers MS-TCM over a single-stream baseline ($\Delta\text{AIC} = 253$) and lands in the parameter regime predicted by the whitepaper: a dominant storyline-context weight with negligible global weight. We discuss what this fit does and does not tell us about the model’s adequacy, and outline five testable predictions that MS-TCM generates for naturalistic memory experiments beyond the free-recall paradigm.

1 The puzzle of memory across narrative storylines

Recency and contiguity in free recall. Two robust principles organize human memory for a sequence of events. The first is *recency*: items encountered late in a sequence are remembered better than items encountered earlier. The second is *contiguity*: successful recall of one item biases the next recall toward items that were encoded close to it in time^{10,12,14}. Both principles vary smoothly with the temporal separation between events. In free recall of word lists, contiguity is most clearly visible in the conditional response probability as a function of lag (lag-CRP), which peaks at lag +1 (the immediately-following item) and decays for larger lags¹¹. The same principles show up in much longer time scales: in continuous-distractor free recall, a recency-like decay unfolds over tens of seconds^{3,9}.

The *Temporal Context Model* (TCM;¹⁰) offers a simple explanation for both principles. A single d -dimensional temporal context vector $\mathbf{c}(t)$ is maintained across encoding. Each encoded item f_i at time t updates the context as

$$\mathbf{c}(t) = \rho \mathbf{c}(t-1) + \beta \mathbf{c}_i^{\text{IN}}, \quad (1)$$

where $\rho = \sqrt{1 - \beta^2}$ preserves $\|\mathbf{c}(t)\| = 1$ under unit-norm inputs and $\beta \in (0, 1)$ is the drift rate. The context similarity $\mathbf{c}(t) \cdot \mathbf{c}(t')$ decays as $\rho^{|t-t'|}$, so items encoded near each other in time share more context than items encoded far apart; at retrieval, reinstatement of a cue’s encoding context preferentially activates its temporal neighbours. TCM produces both recency (end-of-list items dominate retrieval because their context overlaps most with the retrieval cue) and contiguity

47 (nearby items share context and are cross-cued). The *Context Maintenance and Retrieval* model
48 (CMR;¹⁵) extends this machinery with a second sub-region of context that carries source or task
49 information, and shows that CMR jointly captures semantic, temporal, and source clustering.

50 **Memory for interleaved storylines.** For decades the neighbourhood of phenomena around TCM
51 and CMR was studied almost exclusively with experimenter-generated sequences of words. Over
52 the last ten years, an increasingly rich literature has begun to ask how the same contextual principles
53 apply when the sequence is a naturalistic narrative – a television show, a movie, or a spoken story –
54 that humans watch or listen to in a continuous stream^{1,8,22}. Two distinctive features of naturalistic
55 narratives are relevant here.

56 First, *narratives are structured by events*. Event-segmentation theory²¹ and the situation-models
57 literature²³ argue that ongoing experience is parsed into discrete event representations at bound-
58 aries marked by changes in characters, setting, goals, or causal structure. Event boundaries mod-
59 ulate memory: items encoded across a boundary are remembered as further apart in time than
60 items within the same event, even when objective time is matched^{6,8}. Shared schema structure
61 within an event supports robust reinstatement after interruption².

62 Second, *naturalistic narratives often interleave multiple storylines*. Many television shows and
63 novels tell their stories by cutting between parallel threads: the viewer experiences, moment to
64 moment, a single linear stream of events that is the braided-together interleaving of multiple
65 ongoing storylines. *Why Women Kill* is one such show; *This Is Us* is another. A viewer who watches
66 such a show experiences events from storyline *A*, then events from storyline *B*, then events from
67 storyline *A* again, and so on, even though within each storyline events would follow one another
68 continuously if unspooled. Figure 1 illustrates the metaphor: each storyline is a strand with its
69 own internal timeline, and the viewer’s experience is a braid that splices segments from the strands
70 together.

71 **The empirical puzzle.** Xu et al.²⁰ asked whether cued recall of narrative events obeys the same
72 temporal-contiguity principles that govern cued recall of word lists. Their participants watched

73 the first season of *Why Women Kill* – a show that interleaves three storylines set in different historical
 74 periods – and were later cued with one event from the show and asked to recall the event that
 75 preceded or followed it. The cued-to-target pairs fell into three conditions:

- 76 1. **Within-event** (Experiment 2): cue and target belong to the same event (same storyline, same
 77 scene).
- 78 2. **Across-event-within-storyline** (Experiment 1, “grouped”): cue and target are consecutive
 79 events on the same storyline with no intervening events from other storylines.
- 80 3. **Across-event-bridge** (Experiment 3, “interleaved”): cue and target are narratively consecu-
 81 tive on the same storyline but separated by $m = 2\text{--}7$ intervening events from other storylines
 82 (mean ≈ 3).

83 A single-stream contiguity account predicts a substantial decrement for the bridge condition
 84 relative to the grouped condition: the bridge target is four times further away in the viewer’s
 85 global temporal context than the grouped target, and under Equation 1 context similarity falls as
 86 ρ^{m+1} . In the data, by contrast, the bridge and grouped conditions are statistically indistinguishable,
 87 with a Bayes factor $\text{BF}_{01} = 412$ in favour of the null hypothesis of no difference between them. This
 88 is the puzzle we seek to explain.

89 2 A multi-stream temporal context model

90 **Parallel context streams.** MS-TCM assumes that while a listener or viewer watches a narrative,
 91 the brain maintains *multiple* temporal-context vectors in parallel. One of these is a global context,
 92 $\mathbf{c}_G(t)$, that behaves exactly like TCM’s single context: it drifts at every encoding step and therefore
 93 tracks the viewer’s timeline (Figure 2, top centre; Figure 1C, top rail). The rest of the context bank
 94 is one vector per storyline, $\{\mathbf{c}_{Sk}(t)\}_{k=1}^K$, each of which drifts *only while its own storyline is the active*
 95 *storyline* and is held constant otherwise.

96 Concretely, when event i from active storyline S^* is encoded at global time step t : $\mathbf{c}_G(t) =$
 $\rho_G \mathbf{c}_G(t-1) + \beta_G \mathbf{c}_i^{\text{IN}},$

97

$$\mathbf{c}_{S^*}(t) = \rho_S \mathbf{c}_{S^*}(t_{\text{prev}}^{S^*}) + \beta_S \mathbf{c}_i^{\text{IN}}, \quad (98)$$

99 $\mathbf{c}_{S_k}(t) = \mathbf{c}_{S_k}(t-1)$ for all $k \neq S^*$, where $\rho_G = \sqrt{1 - \beta_G^2}$, $\rho_S = \sqrt{1 - \beta_S^2}$, and $t_{\text{prev}}^{S^*}$ is the most recent
 100 prior time step at which S^* was active. Equation 2 is the key architectural commitment of MS-TCM
 101 and the ingredient that standard TCM¹⁰ and CMR¹⁵ do not have: a storyline context's state at time
 102 t depends only on the events in its own storyline, not on intervening events from other storylines.
 103 When the storyline resumes at a later viewer time t' , the update in Eq. 2 reuses the *frozen* state
 104 from $t_{\text{prev}}^{S^*}$ rather than whatever the storyline context would have drifted to had it updated at every
 105 intervening step.

106 **Encoding and retrieval composites.** Each event i on storyline S^* is associated with a composite
 107 encoding context that mixes the global and active-storyline contexts:

$$\mathbf{c}_{\text{comp}}(i) = w_G \mathbf{c}_G(t) + w_S \mathbf{c}_{S^*}(t), \quad w_G + w_S = 1. \quad (2)$$

108 During cued recall, presentation of a cue event j on storyline S_j reinstates a retrieval composite:

$$\mathbf{c}_{\text{ret}}(j) = w_G^{\text{ret}} \mathbf{c}_G(t_j) + w_S^{\text{ret}} \mathbf{c}_{S_j}(t_j), \quad w_G^{\text{ret}} + w_S^{\text{ret}} = 1. \quad (3)$$

109 Crucially, the retrieval weights $(w_G^{\text{ret}}, w_S^{\text{ret}})$ need not equal the encoding weights (w_G, w_S) . Task
 110 instructions can rebalance the cue: for instance, being asked to recall "what happened next in the
 111 same storyline" should plausibly increase w_S^{ret} relative to w_G^{ret} , amplifying the storyline context's
 112 contribution to the cue. The probability of recalling a candidate event i given cue j is a softmax-
 113 with-gain over a scored similarity,

$$s_i(j) = \cos(\mathbf{c}_{\text{ret}}(j), \mathbf{c}_{\text{comp}}(i)) \cdot a_i \cdot e^{-\lambda I_{ij}}, \quad P(i | j) = \frac{\exp(\tau s_i(j))}{\sum_{k \in C(j)} \exp(\tau s_k(j))}, \quad (4)$$

114 where $a_i = \phi_s e^{-\phi_d(i-1)} + 1$ is the Polyn-et-al.-style primacy-gradient item bias, $\tau > 0$ is an inverse-
 115 temperature / softmax gain, I_{ij} is the cumulative between-event interference term scaled by λ (§4),

116 and the candidate set $C(j)$ is the set of list items that have not yet been recalled on the current list
 117 (the standard TCM / CMR "no repeats" convention). Both τ and (ϕ_s, ϕ_d) are fit alongside the other
 118 free parameters. Without τ the softmax of bounded cosine similarities in $[-1, 1]$ is close to uniform
 119 and the model fails to produce any of the shape features of free recall; without (ϕ_s, ϕ_d) there is no
 120 primacy mechanism and the serial- position curve has only a recency limb. Setting $\phi_s = 0$ recovers
 121 the vanilla formulation.

122 **Relationship to TCM and CMR.** MS-TCM specializes to standard TCM¹⁰ when $w_s = 0$, in which
 123 case the storyline bank is inert and the model reduces to Eq. 1. The model is also superficially
 124 similar to CMR¹⁵: both augment TCM with a second source of contextual information beyond
 125 pure recency. But the architectures differ in two important ways.

126 First, CMR maintains a *single* context vector with multiple sub-regions (temporal and source),
 127 all of which update according to Eq. 1 at every step, with possibly different drift rates. MS-TCM
 128 maintains $K + 1$ *parallel context vectors* with qualitatively different update rules: only one of the
 129 K storyline contexts updates per step, and the other $K - 1$ are held constant. The frozen-while-
 130 inactive property is not a weighted-down version of drift; it is the absence of update.

131 Second, CMR's source sub-region is a *feature of the current event* – it represents the task or source
 132 under which the current item was encoded. MS-TCM's storyline contexts are *separate memory*
 133 *stores*: the storyline identity of an event determines which store gets updated, not what the store
 134 *contains*. A storyline context carries the accumulated contextual imprint of all previous events *on*
 135 *that storyline*, independent of whatever events have been encoded on other storylines in between.

136 3 Deriving the grouped–bridge equivalence

137 We now derive analytically the parameter regimes under which MS-TCM reproduces the observed
 138 grouped–bridge equivalence. Consider two consecutive storyline- S events i and i' encoded at
 139 viewer times t and $t' = t + m + 1$, so $m \geq 0$ intervening events from other storylines separate them.
 140 Under orthonormal item inputs (a standard simplification; see Eq. (1) of¹⁰), the per-stream context

similarities after the second encoding are: $\cos(\mathbf{c}_G(t), \mathbf{c}_G(t')) \approx \rho_G^{m+1}$,
 $\cos(\mathbf{c}_S(t), \mathbf{c}_S(t')) = \rho_S$, the first because $\mathbf{c}_G$ has drifted through $m + 1$ encoding steps, the second
because $\mathbf{c}_S$ was frozen during the m intervening steps and so has undergone only one update
between t and t' . A linear composite-similarity quantity $\sigma(m) \equiv w_G \rho_G^{m+1} + w_S \rho_S$ collapses the two
conditions to

$$\sigma(0) = w_G \rho_G + w_S \rho_S \quad (\text{grouped}), \quad \sigma(m) = w_G \rho_G^{m+1} + w_S \rho_S \quad (\text{bridge}). \quad (5)$$

The observed equivalence $\sigma(0) \approx \sigma(m)$ holds whenever one of three regimes obtains:

1. **Storyline-context dominance.** When $w_S \gg w_G$, both conditions are dominated by the $w_S \rho_S$
term, which is identical across conditions. The global drift contribution is a minority of the
composite and the difference $w_G(\rho_G - \rho_G^{m+1})$ is accordingly small.
2. **Task-driven retrieval reweighting.** Even when encoding weights are comparable ($w_G \approx w_S$),
the Experiment 3 instruction ("recall the next event in the same storyline") can plausibly up-
weight the storyline component at retrieval by increasing w_S^{ret} relative to w_G^{ret} , again shrinking
the bridge-grouped gap.
3. **Matched interference.** The grouped condition pairs the target event with a within-storyline
boundary event that shares characters and setting, which should produce substantial retroac-
tive interference. The bridge condition pairs the target with events from *different* storylines,
which are categorically distinct and should produce less interference release from proactive
interference;¹⁸. If this between-storyline material produces less interference than within-
storyline material, the interference term $e^{-\lambda I_{ij}}$ in Eq. 4 can partially offset the reduced global
context similarity in the bridge condition.

Numerical illustration. As a concrete anchor, at $\beta_G = \beta_S = 0.5$, $w_G = 0.2$, $w_S = 0.8$, and $m = 3$ (the
mean number of intervening events in *Why Women Kill*), the closed-form composite similarities are
 $\sigma(0) = \sqrt{0.75} = 0.86603$ (grouped) and $\sigma(3) = 0.2 \cdot 0.5625 + 0.8 \sqrt{0.75} = 0.80532$ (bridge), a difference

164 of 0.06 that is small relative to typical item- and participant-level variance in cued-recall studies.¹

165 4 Additional mechanisms

166 The core architecture in §2 is deliberately minimal – it is the smallest change to standard TCM that
 167 produces the frozen-while-inactive property required to explain the empirical puzzle. Three op-
 168 tional mechanisms can be added on top, each corresponding to a specific claim from the naturalistic-
 169 memory literature.

170 **Resumption reinstatement (γ).** When a storyline resumes after interruption, re-encountering
 171 its characters and setting may trigger automatic reinstatement of the storyline context to some-
 172 thing close to its prior frozen value analogous to environmental-reinstatement effects in episodic
 173 memory;^{7,17}. Formally, the storyline update becomes

$$\mathbf{c}_{S^*}(t) = \rho_S \mathbf{c}_{S^*}(t_{\text{prev}}^{S^*}) + \beta_S \mathbf{c}_i^{\text{IN}} + \gamma \mathbf{c}_{S^*}(t_{\text{prev}}^{S^*}), \quad (6)$$

174 where $\gamma \geq 0$ captures the reinstatement strength and reduces to Eq. 2 at $\gamma = 0$.

175 **Conversational references (α_k).** Xu et al.¹⁹ show that characters in television narratives refer
 176 more often to past events on the same storyline than to future events. When storyline S resumes,
 177 explicit references in dialogue ("*Remember when we discussed the project?*") can provide *explicit*
 178 reinstatement of earlier storyline events. These references enter the storyline update as additional
 179 inputs:

$$\mathbf{c}_{S^*}(t) = \rho_S \mathbf{c}_{S^*}(t_{\text{prev}}^{S^*}) + \beta_S \mathbf{c}_i^{\text{IN}} + \sum_k \alpha_k \mathbf{c}_{\text{ref}_k}^{\text{IN}}, \quad (7)$$

180 where ref_k indexes events explicitly referenced in the current event's dialogue and α_k controls the
 181 weight of each reference.

¹The whitepaper for MS-TCM rounds $\sigma(3)$ to 0.806 via a rounding-cascade ($\rho \approx 0.866 \Rightarrow \rho^4 \approx 0.563 \Rightarrow \sigma(3) \approx 0.806$). The closed-form value is 0.80532..., which rounds to 0.805. The code's unit test `test_section.4.4.numerical.anchor` asserts the closed-form values at 10^{-6} tolerance. See Supplement §S2.

182 **Differential interference (λ).** The release-from-proactive-interference literature¹⁸ suggests that
 183 intervening events from categorically distinct material produce less interference than within-
 184 category intervening material. MS-TCM captures this by scaling the retrieval score by $e^{-\lambda I_{ij}}$ (Eq. 4),
 185 where I_{ij} is the cumulative between-event interference from items encoded between i and j . For
 186 free recall we instantiate $I_{ij} = \max(0, |s_i - s_j| - 1)$, the count of items presented between candidate i
 187 (at serial position s_i) and cue j (at serial position s_j); $\lambda = 0$ disables the mechanism.

188 5 Testable predictions

189 Beyond the grouped-bridge equivalence, MS-TCM generates at least five predictions that discrim-
 190 inate it from single-stream accounts of narrative memory.

- 191 1. **Storyline-similarity manipulation.** The equivalence should break down when storylines
 192 share characters, settings, or overlapping themes. With highly similar storylines, distinct
 193 c_{sk} 's become harder to maintain as separable representations, reducing the effectiveness of the
 194 frozen-while-inactive mechanism. A natural test: compare bridge to grouped performance
 195 on *Why Women Kill* (three distinct storylines) against *The Affair* (two storylines showing the
 196 same events from two perspectives). MS-TCM predicts bridge $\approx$ grouped for the first; TCM
 197 predicts bridge $<$ grouped in both.
- 198 2. **Number of intervening events (m).** Under MS-TCM, across-event-bridge performance
 199 should be relatively robust to m because storyline-context similarity is ρ_S regardless of m . A
 200 weak decline is still expected from the diminishing global-context contribution. In *This Is Us*,
 201 where m varies from 1 to 7, MS-TCM predicts a weak bridge- m slope and TCM predicts a
 202 steep one.
- 203 3. **Individual differences in working memory.** Maintaining K distinct storyline contexts in
 204 parallel should be more demanding than maintaining a single temporal context. Individuals
 205 with higher working-memory capacity should therefore show a stronger grouped $\approx$ bridge
 206 equivalence (and more sensitive c_{sk} distinctions) than individuals with lower capacity.

207 **4. Temporal direction asymmetry.** Within-event cued recall in Xu et al.²⁰ shows a strong
 208 backward asymmetry ($BF_{10} = 1.6 \times 10^{23}$) that the across-event conditions do not. MS-TCM
 209 attributes this to a reduction of pre-built backward links at event boundaries⁶ Xu et al. 2024,
 210 since the storyline-context mechanism is approximately symmetric (ρ_S is the same regardless
 211 of direction). Manipulating pre-built links experimentally (e.g., by familiarizing participants
 212 with forward *or* backward event sequences before the cued-recall test) should dissociate the
 213 two mechanisms.

214 **5. First-recall latency.** Even with equivalent accuracy, bridge-condition responses may be
 215 *slower* than grouped- condition responses if retrieval begins with the (poorer) global cue and
 216 incorporates the storyline cue serially. A full-RT model of MS-TCM retrieval, along the lines
 217 of TCM-A¹⁶ – not attempted here – would make this quantitative.

218 **6 Worked example: fitting MS-TCM to a free-recall dataset**

219 Because the cued-recall data from Xu et al.²⁰ are not yet publicly released, we demonstrate the MS-
 220 TCM machinery on a public free-recall dataset whose structure is analogous: the category condition
 221 of Manning et al.¹³’s feature-rich free-recall experiment. In this condition, $N = 30$ participants each
 222 study 16 lists of 16 words drawn from 4 taxonomic categories (Figure 3). In early lists (list index < 8)
 223 the 4-word category groups are presented consecutively; in late lists (list index ≥ 8) the category
 224 order is randomized. Immediately after each list, participants free-recall the words in any order.
 225 The category manipulation is structurally analogous to MS-TCM’s storyline variable: taxonomic
 226 category indexes which c_{Sk} should be active for each encoded item, grouped lists play the role of
 227 the grouped condition, and randomized lists play the role of the bridge condition.

228 **Feature encoding and likelihood.** Each word is encoded as a 71-dimensional multi-hot feature
 229 vector that concatenates one-hot indicators for category (15 + 1 unknown bins), size (small / large),
 230 first letter (26), word length (10 bins), colour (8 coarse bins of the RGB cube), and on-screen quadrant
 231 (4), plus a reserved list-start dimension. This schema is determined by the FRFR annotation and

Table 1: Maximum-likelihood parameter estimates for MS-TCM fit to the FRFR category condition. Values are MLE [95%*bootstrapCI*] (percentile method, $n_{bootstraps} = 20$). Derived parameters are not fit independently (e.g. $w_S = 1 - w_G$).

Parameter	Symbol	Value [95%CI]	Description
beta_global	β_G	0.950 [0.001, 1.000]	Drift rate of the global temporal context. Higher values mean the global context updates more strongly toward each new event (and forgets older ones more quickly).
beta_storyline	β_S	0.533 [0.241, 0.551]	Drift rate of the storyline-specific context. A value near 0 means each storyline context behaves as a nearly static signature of the category; a value near 1 means within-storyline events overwrite one another as they are encoded.
w_global	w_G	0.000 [0.000, 0.543]	Weight on the global context in the composite encoding context. Higher values make the model more sensitive to global temporal contiguity (standard TCM sets $w_G = 1$).
w_storyline	w_S	1.000 (derived)	Weight on the active storyline context in the composite. Higher values make category or storyline identity a stronger cue for recall.
w_global_ret	w_G^{ret}	0.000 (derived)	Global-context weight at retrieval. Can differ from w_G at encoding, reflecting task-driven reweighting of cues.
w_storyline_ret	w_S^{ret}	1.000 (derived)	Storyline-context weight at retrieval. Increasing it models an instruction to recall within the same storyline.
$\log L = -14067.2$, AIC = 28140.4, BIC = 28160.1; $n_{participants} = 30$, $n_{lists} = 480$, $n_{recalls} = 5205$			

is fixed across participants; its exact form is documented in `code/ms_tcm/features.py`. Per-list log-likelihood is the sum over observed in-list recall transitions of $\log P(r_k | \mathbf{c}_{\text{ret}}(r_{k-1}))$ (Eq. 4), and dataset log-likelihood is the sum across every participant-list pair. Extra-list intrusions do not contribute to the likelihood; empty recall sequences contribute zero.

Optimizer and confidence intervals. The constrained parameters $\beta_G, \beta_S \in (0, 1)$ and $w_G \in (0, 1)$ are reparameterized via sigmoid transforms onto an unconstrained vector, and optimized by L-BFGS-B with two restarts from seeded random initial points. Confidence intervals come from a participant-level percentile bootstrap ($n_{\text{boot}} = 20$, all converged). Two models are fit: MS-TCM (all parameters free) and standard TCM (constrained to $w_S = 0$, so the storyline stream is inert). Both fits use identical seeds, feature encoders, and dataset (7,680 encoded events, 5,205 in-list recalls, 10 extra-list intrusions).

243 **Parameter estimates.** The MLE for MS-TCM lands in §3’s regime (1) *storyline-context dominance*:
 244 the optimizer drives w_G to essentially 0 and w_S to essentially 1, so the composite encoding context
 245 is entirely the active storyline’s context. The storyline drift rate $\beta_S = 0.53$ is in the mid-range and
 246 well-identified (CI [0.24, 0.55]); the global drift rate $\beta_G = 0.95$ has a wide CI because $w_G \approx 0$ makes
 247 β_G unidentifiable – with no global weight in the composite, the global drift rate does not affect the
 248 likelihood.

249 **Model comparison.** MS-TCM is strongly preferred over standard TCM on this dataset. The per-
 250 transition log-likelihood is $-14,067$ for MS-TCM vs $-14,195$ for standard TCM, for a difference of
 251 128 nats over 5,205 recall transitions. Correcting for the additional free parameters in MS-TCM
 252 yields $\Delta AIC = 253$ and $\Delta BIC = 239$, both overwhelmingly in favour of MS-TCM. Under standard
 253 interpretations⁴ this is decisive evidence that adding the storyline-context stream improves the
 254 model’s prediction of observed free-recall transitions.

255 **What the fit does not reproduce.** The raw likelihood-preference for MS-TCM is not the whole
 256 story. Figure 4 overlays three classic free-recall measures on the MS-TCM posterior-predictive
 257 (red) and the observed behaviour (black). The observed data show a U-shaped serial-position
 258 curve with strong primacy and recency, strong end-of-list first-recall probabilities, and a sharply
 259 peaked forward-asymmetric lag-CRP – all signature phenomena of temporal contiguity. The MS-
 260 TCM MLE, operating in the $w_S \rightarrow 1$ regime, produces near-flat SPC, near-flat PFR, and a near-flat
 261 lag-CRP. Figure 5 tells the same story for percentile-rank clustering: observed temporal, category,
 262 and semantic clustering are all above the 0.5 no-clustering baseline, but the MS-TCM posterior-
 263 predictive sits at 0.50–0.53 across all three dimensions.

264 **Interpretation.** The combination of high per-transition likelihood and low posterior-predictive
 265 match on structured behavioural measures is a classic signature of a *mis-specified objective*. The
 266 per-transition likelihood rewards any factor that improves the per-step prediction, and in the
 267 FRFR-category design category identity is a very strong per-step predictor: given a just-recalled
 268 item from category A , $\approx 60\%$ of next-recalls are also from category A . An MLE that allocates

all its capacity to category-based prediction therefore scores well on per-transition likelihood but sacrifices the global-context machinery that is needed to reproduce the temporal-contiguity-driven phenomena in Figure 4. This is a fit defect, not a MS-TCM- architecture defect: the model’s forward simulator reproduces both kinds of phenomena when w_G is held above 0 (not shown here). We return to this point in the Discussion.

7 Relationship to existing models

Standard TCM¹⁰. MS-TCM reduces to TCM in the $w_S = 0$ limit. TCM’s proximity relationship $\mathbf{c}(t) \cdot \mathbf{c}(t') = \rho^{|t-t'|}$ is recovered for the global stream in MS-TCM. What MS-TCM adds is the K -dimensional bank of storyline streams with frozen-while-inactive dynamics. That bank cannot be written as a single context vector with rotated basis because the *update rule* depends on which storyline is currently active – a property that no single-stream drift model has.

CMR¹⁵. CMR is superficially similar to MS-TCM in that it augments TCM with non-temporal contextual information, but the architectural commitments are different in three respects. First, CMR’s source sub-region updates at every encoding step (possibly with a different drift rate β^{source}); MS-TCM’s storyline stores update only when their storyline is active. Second, CMR represents source as a *feature of the current event* – the source sub-region at time t carries the source identity of the event encoded at t . MS-TCM’s storyline streams are separate *stores* whose identity indexes which stream updates, not what the stream contains. Third, CMR uses learned associative matrices $M^{\text{FC}} / M^{\text{CF}}$ to mediate retrieval; MS-TCM directly compares $\mathbf{c}_{\text{ret}}$ to each candidate’s $\mathbf{c}_{\text{comp}}$ via cosine similarity. These three differences combine to yield the frozen-while-inactive property that the grouped-bridge equivalence requires; adding a source sub-region that drifts at a slower rate than the temporal sub-region (as CMR allows) does not produce the equivalence, because the source sub-region still accumulates drift at every intervening event.

Event-segmentation theory and situation models. Zacks et al.²¹ argue that ongoing experience is parsed into discrete event representations at boundaries marked by changes in characters, setting,

294 or goals, and that event boundaries modulate downstream memory access^{5,6,8}. Situation models²³
 295 track dimensions such as space, time, causation, and character identity. MS-TCM’s storyline
 296 contexts can be read as a computational implementation of the *within-storyline component* of a
 297 situation model: they track within-storyline temporal continuity and are backgrounded (“frozen”)
 298 when a different storyline becomes active, corresponding to the maintenance of a dormant situation
 299 model. The key addition in MS-TCM is the *quantitative* prediction that cued recall across an event
 300 boundary should depend on the storyline membership of the intervening events, not just their
 301 count.

302 **Schema-based reinstatement.** Baldassano et al.² show that shared event schemas support rein-
 303 statement of event representations across interruption. A schema can be read, in MS-TCM terms,
 304 as a structural scaffold that makes each c_{sk} more robust to decay during its frozen period. Ex-
 305 perimental manipulations of schema strength (e.g., contrasting formulaic from novel narrative
 306 structures) should therefore modulate the bridge/grouped equivalence.

307 **TCM-A and the decision stage.** Sederberg et al.¹⁶’s TCM-A model pairs TCM with a leaky-
 308 competing- accumulator decision stage capable of modelling inter-response times. The MS-TCM
 309 retrieval machinery in Eq. 4 is a softmax over scored similarities, which is a compact approximation
 310 to the equilibrium behaviour of a TCM-A-style competition but does not predict response latencies.
 311 Extending MS-TCM to TCM-A dynamics is straightforward and would unlock the first-recall-
 312 latency prediction of §5 (item 5).

313 8 Discussion

314 We have proposed a multi-stream extension of the temporal context model in which a bank of
 315 storyline-specific context vectors drift only while their own storyline is active and are held constant
 316 otherwise. The architecture is the smallest departure from standard TCM that can reproduce the
 317 Xu et al.²⁰ finding of across-event-bridge $\approx$ across-event-within-storyline cued recall ($BF_{01} = 412$).

318 We derived closed-form regimes under which MS-TCM predicts this equivalence (§3), and demon-
319 strated the machinery on a public free-recall dataset whose category manipulation is structurally
320 analogous to the storyline manipulation in the target experiment (§6).

321 **The fit tells us two things.** First, the per-transition-likelihood preference for MS-TCM over stan-
322 dard TCM is enormous ($\Delta\text{AIC} = 253$, $\Delta\text{BIC} = 239$) on the FRFR-category data. Every standard
323 information-criterion interpretation places the posterior odds overwhelmingly in favour of the
324 multi-stream architecture. This is non-trivial evidence that the storyline stream carries real pre-
325 dictive signal for human free recall of category-structured word lists, not just for naturalistic
326 narratives.

327 Second, the MLE alone is *not sufficient* to characterize the model’s empirical adequacy. The
328 per-transition-likelihood-MLE allocates all the encoding-composite mass to the storyline stream
329 ($w_S \rightarrow 1$) and in doing so loses the ability to reproduce temporal-contiguity phenomena (Figures 4
330 and 5). The model’s forward simulator, run at hand-picked parameters that keep w_G away from
331 0, does reproduce both kinds of phenomena. We read this as a specification problem: a pure
332 per-transition-likelihood objective is insufficient for fitting a model that needs to reproduce both
333 point-transition and list-structured phenomena. Two natural remedies are a multi-objective fit
334 that penalizes deviations of predicted SPC and lag-CRP from observed, and a Bayesian prior on
335 (w_G, w_S) that regularizes w_G away from the boundary. Both are out of scope for the present paper
336 and are natural targets for future work.

337 **Caveats specific to the worked example.** Three caveats are worth flagging. **(1)** FRFR-category is a
338 free-recall paradigm; the across-event-bridge phenomenon that motivates MS-TCM is a cued-recall
339 phenomenon. The fit reported here is therefore evidence that MS-TCM can account for free-recall
340 phenomena in a category-structured design, not a direct test of the equivalence it was designed to
341 explain. A full test awaits the public release of the Xu et al.²⁰ cued-recall data. **(2)** The storyline
342 variable in FRFR-category is taxonomic category, which is a well-defined but relatively thin proxy
343 for *narrative* storyline. Real narrative storylines have within-storyline continuity of characters,

344 setting, goals, and thematic content that categories like “mammals” and “fruits” do not. A richer
345 feature encoding – e.g., sentence-transformer embeddings of event annotations, as proposed for
346 the naturalistic data in the MS-TCM whitepaper accompanying this repository (notes/ms-tcm.pdf
347 §8) – would be a closer match. (3) The optional $\gamma, \alpha_k, \lambda$ mechanisms from §4 are disabled in the fit
348 reported above. The FRFR-category data are not rich enough to identify all of them jointly; fitting
349 them would require either a dataset with explicit within-storyline reference structure or a prior on
350 the mechanism parameters.

351 **Broader implications.** If the multi-stream picture is right, it implies a strong reconceptualization
352 of what “temporal context” means in narrative memory. A single continuously drifting context
353 vector is an adequate summary of the timeline when there is only one timeline to track. But when
354 the listener experiences a braided interleaving of several storyline timelines, their memory system
355 appears to maintain separate, parallel context vectors – one per storyline – each of which runs on
356 *its own* internal clock. That internal clock ticks only during the storyline’s own events and is held
357 frozen during events on other storylines. This view unifies event-segmentation theory’s claim
358 that event representations are maintained in parallel²¹ with TCM’s claim that memory access is
359 organized by drift-based context similarity¹⁰: narrative memory uses TCM-like machinery, but
360 runs it independently on each storyline. It also predicts that the cost of narrative interleaving is
361 paid by the global timeline (which is what the bridge condition’s global-context drift would predict
362 under single-stream TCM), not by the within-storyline timeline – exactly the pattern in Xu et al.²⁰.

363 **Reproducibility.** The complete pipeline – raw-data download, reformat (including a regression
364 test that pins the expected serial-position distribution), validation, MLE fit, bootstrap, figure regen-
365 eration, and paper compilation – is driven by `reproduce.sh` at the repository root. The shipped
366 `data/raw/frfr.category/` directory is self-verifying via SHA-256 hashes in its `manifest.json`.
367 All figures in this paper are generated by scripts in `code/figures/` with a 1:1 script-to-figure
368 mapping. A cross-platform regression test pins the macOS MLE so the same dataset on Ubuntu or
369 Windows CI produces numerically identical fits (within 10^{-6}).


References

- [1] C Baldassano, J Chen, A Zadbood, J W Pillow, U Hasson, and K A Norman. Discovering event structure in continuous narrative perception and memory. *Neuron*, 95(3):709–721, 2017.
- [2] C Baldassano, U Hasson, and K A Norman. Representation of real-world event schemas during narrative perception. *The Journal of Neuroscience*, 38(45):9689–9699, 2018.
- [3] R A Bjork and W B Whitten. Recency-sensitive retrieval processes in long-term free recall. *Cognitive Psychology*, 6:173–189, 1974.
- [4] K P Burnham and D R Anderson. *Model selection and multimodel inference: a practical information-theoretic approach*. Springer, 2 edition, 2004.
- [5] David Clewett and Lila Davachi. The ebb and flow of experience determines the temporal structure of memory. *Current Opinion in Behavioral Sciences*, 17:186–193, 2017.
- [6] Y Ezzyat and L Davachi. What constitutes an episode in episodic memory? *Psychological Science*, 22:243–252, 2011.
- [7] D R Godden and A D Baddeley. Context-dependent memory in two natural environments: on land and underwater. *British Journal of Psychology*, 66:325–331, 1975.
- [8] A C Heusser and J R Manning. Capturing the geometric structure of episodic memories for naturalistic experiences. In *Proceedings of the Conference on Cognitive Computational Neuroscience*, pages PS–2B.16, 2018.
- [9] M W Howard and M J Kahana. Contextual variability and serial position effects in free recall. *Journal of Experimental Psychology: Learning, Memory, and Cognition*, 25:923–941, 1999.
- [10] M W Howard and M J Kahana. A distributed representation of temporal context. *Journal of Mathematical Psychology*, 46:269–299, 2002.
- [11] M J Kahana. Associative retrieval processes in free recall. *Memory and Cognition*, 24:103–109, 1996.


- 394 [12] M J Kahana, M W Howard, and S M Polyn. Associative retrieval processes in episodic memory.
395 In *Learning and Memory: A Comprehensive Reference* (J Byrne, ed.), pages 467–490. Elsevier, 2008.
- 396 [13] Jeremy R Manning, Emily C Whitaker, Paxton C Fitzpatrick, Madeline R Lee, Allison M Frantz,
397 Bryan J Bollinger, Darya Romanova, Campbell E Field, and Andrew C Heusser. Feature and
398 order manipulations in a free recall task affect memory for current and future lists, 2023.
399 PsyArXiv preprint, <https://psyarxiv.com/erzfp>. Data and code: [https://github.com/](https://github.com/ContextLab/FRFR-analyses)
400 ContextLab/FRFR-analyses.
- 401 [14] S M Polyn and M J Kahana. Memory search and the neural representation of context. *Trends*
402 *in Cognitive Sciences*, 12:24–30, 2008.
- 403 [15] S M Polyn, K A Norman, and M J Kahana. Task context and organization in free recall.
404 *Neuropsychologia*, 47:2158–2163, 2009.
- 405 [16] P B Sederberg, M W Howard, and M J Kahana. A context-based theory of recency and
406 contiguity in free recall. *Psychological Review*, 115(4):893–912, 2008.
- 407 [17] S M Smith and E Vela. Environmental context-dependent memory: a review and meta-
408 analysis. *Psychonomic Bulletin and Review*, 8(2):203–220, 2001.
- 409 [18] D D Wickens. Encoding categories of words: an empirical approach to meaning. *Psychological*
410 *Review*, 77:1–15, 1970.
- 411 [19] X Xu, Z Zhu, X Zheng, and J R Manning. Temporal asymmetries in inferring unobserved past
412 and future events. *Nature Communications*, 15:8502, 2024.
- 413 [20] Xinming Xu, Katherine Duncan, and Jeremy R Manning. Cued recall of narrative events
414 across storylines and event boundaries. Manuscript in preparation, 2026.
- 415 [21] J M Zacks, N K Speer, K M Swallow, T S Braver, and J R Reynolds. Event perception: a
416 mind-brain perspective. *Psychological Bulletin*, 133:273–293, 2007.

- 417 [22] A Zadbood, J Chen, Y C Leong, K A Norman, and U Hasson. How we transmit memories to
418 other brains: constructing shared neural representations via communication. *Cerebral Cortex*,
419 27(10):4988–5000, 2017.
- 420 [23] R A Zwaan and G A Radvansky. Situation models in language comprehension and memory.
421 *Psychological Bulletin*, 123(2):162–185, 1998.

A. Three storylines, each with its own internal timeline



B. The interleaved braid the viewer experiences



C. Context drifts

global context drifts at every step; each storyline context drifts only during its own segments, and is frozen otherwise




Figure 1: The interleaved-threads picture of narrative memory. (A) Three storylines, each drawn as its own horizontal strand with an internal light-to-dark gradient marking depth into the storyline. Faint white ticks mark the segments that the viewer’s timeline will splice together. (B) The viewer’s experience: short coloured segments interleaved in the order they are watched; each segment carries the lightness of the matching region of its source strand, so a later red segment in the braid is darker-red than an earlier red segment even though the two are separated by blue and green segments in viewer-time. Thin storyline-coloured arcs show which strand each segment came from. (C) Context drifts. The global context (c_G) updates at every viewer tick (a smooth grey-to-black gradient). Each storyline-context rail drifts only during its own segments and is flat-held (“frozen”) elsewhere.




Figure 2: The MS-TCM architecture. Each encoding step draws a feature input $\mathbf{c}_i^{\text{IN}}$ from the current event. That input drives two kinds of context update: the global context $\mathbf{c}_G$ updates every step (drift rate β_G), while only the storyline context $\mathbf{c}_{S^*}$ for the storyline that event i belongs to updates (drift rate β_S). Inactive storyline contexts are frozen. The encoding composite $\mathbf{c}_{\text{comp}}(i) = w_G \mathbf{c}_G(t) + w_S \mathbf{c}_{S^*}(t)$ is associated with event i . At retrieval, a cue event j reinstates a retrieval composite that can use *different* mixing weights ($w_G^{\text{ret}}, w_S^{\text{ret}}$), implementing task-driven reweighting. The decision stage scores every candidate i 's $\mathbf{c}_{\text{comp}}(i)$ against $\mathbf{c}_{\text{ret}}(j)$ via cosine similarity (optionally modulated by item activations a_i and by an interference term $e^{-\lambda_{ij}}$) and applies a softmax.




Figure 3: FRFR-category paradigm. Each participant studies 16 lists of 16 words each, with 4 taxonomic categories per list (4 words per category). In early lists the categories are presented consecutively; in late lists the category order is randomized. Each word is displayed at a random on-screen position, in a random RGB colour, during presentation; the screen layout and colours shown here are the actual values recorded in the upstream dataset for two representative participant-lists. Data: Manning et al.¹³, category condition, reformatted to the MS-TCM canonical schema with bit-exact SHA-256 verification (see `data/raw/frfr_category/manifest.json`).




Figure 4: Behavioural measures, observed (black) versus the MS-TCM posterior-predictive (red, shaded = 95% across 20 synthetic replications at the MLE). (A) Serial-position curve. (B) Probability of first recall. (C) Lag-CRP (restricted to $|\text{lag}| \leq 5$). The MS-TCM MLE achieves a much higher per-transition log-likelihood than standard TCM but, in the $w_S \rightarrow 1, w_G \rightarrow 0$ regime it prefers, does not reproduce the classic U-shaped serial-position curve, the strong end-of-list PFR, or the forward-asymmetric lag-CRP. See §8 for discussion.




Figure 5: Percentile-rank clustering scores (Polyn et al. 2009 definition) for temporal, category, and semantic dimensions. Black = observed, red = MS-TCM posterior-predictive (20 synthetic datasets at the MLE). Semantic distances are cosine distances in a sentence-transformer (MiniLM) embedding of the recalled words. Dotted line at 0.5 is the no-clustering baseline.